

nlu_rag_eda_synth

November 11, 2025

[]:

1 NLU and RAG Templates EDA

This notebook inspects the synthetic NLU dataset (user queries) and the RAG template corpus used for retrieval. It includes tokenization examples, label distributions, synonym augmentation, and simple template matching checks.

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
nlu = pd.read_csv('./nlu_dataset.csv')
print('NLU data sample:')
nlu.head()
```

NLU data sample:


```
[2]:          user_query    problem_type    domain \
0  Create a voice activity detector for audio rec...  classification  tabular
1  Create a voice activity detector for audio rec...  classification  tabular
2  Create a voice activity detector for audio rec...  classification  tabular
3  Cluster customers into segments based on purch...    clustering   tabular
4  Build a supervised classification model to pre...  classification  tabular

  target_variable  algorithm_preference
0              NaN                  NaN
1              NaN                  NaN
2              NaN                  NaN
3              NaN                  NaN
4        churn                  NaN
```

1.1 Label distribution

```
[3]: print(nlu['problem_type'].value_counts())
nlu['problem_type'].value_counts().plot(kind='bar')
plt.title('Problem type distribution')
plt.show()
```

```
problem_type
classification    93
nlp                52
regression        37
clustering         18
Name: count, dtype: int64
```



1.2 Tokenization and keyword density

```
[4]: import re
nlu['token_count'] = nlu['user_query'].apply(lambda x: len(re.findall(r"\w+", x)))
keywords = [
    'predict', 'classify', 'regress', 'image', 'text', 'cluster', 'summarize', 'detect', 'forecast']
def keyword_density(s):
    s2 = s.lower()
```

```

    return sum(s2.count(k) for k in keywords)/max(1,len(s2.split())))
nlu['keyword_density'] = nlu['user_query'].apply(keyword_density)
nlu[['token_count','keyword_density']].describe()

```

[4]:

	token_count	keyword_density
count	200.000000	200.000000
mean	9.055000	0.121661
std	2.758627	0.081902
min	6.000000	0.000000
25%	7.000000	0.076923
50%	8.000000	0.125000
75%	12.000000	0.166667
max	15.000000	0.285714

1.3 RAG templates preview

[6]:

```

import json
templates = []
with open('./rag_templates.jsonl','r') as f:
    for line in f:
        templates.append(json.loads(line))
print('Templates loaded:', len(templates))
for t in templates:
    print('-', t['template_id'], t['task_type'], t['framework'])

```

Templates loaded: 5

- tpl_0 classification scikit-learn
- tpl_1 regression scikit-learn
- tpl_2 nlp sklearn
- tpl_3 vision pytorch
- tpl_4 time-series statsmodels

1.4 Simple matching test (keyword -> template task)

[7]:

```

def match_template(query):
    q = query.lower()
    if 'image' in q or 'cat' in q or 'dog' in q:
        return 'vision'
    if 'price' in q or 'forecast' in q or 'regress' in q:
        return 'regression'
    if 'review' in q or 'text' in q or 'sentiment' in q:
        return 'nlp'
    return 'classification'

nlu['matched_task'] = nlu['user_query'].apply(match_template)
import pandas as pd
print(pd.crosstab(nlu['problem_type'], nlu['matched_task']))

```

matched_task	classification	nlp	regression	vision
problem_type				
classification	28	16	5	44
clustering	18	0	0	0
nlp	39	5	8	0
regression	10	8	19	0

1.5 Observations

- NLU intents are extractable with keyword heuristics, but an LLM-based NLU will improve robustness.
- RAG templates are small; for production, enlarge corpus and index embeddings.

[]:

[]: